



Eight fixed held-out tasks: paired differences and 95% bootstrap intervals



64-token cap: baseline 8/8 and selected 6/8 outputs reach the cap.
Conditional on these tasks; answer inclusion is not a complete quality assessment.